

The Architecture of Symbolic Computers

Peter M. Kogge

Preface

Part I Fundamentals

Chapter 1 Fundamentals of Computation

- 1.1 Problems and Algorithms
- 1.2 Languages: Syntax and Semantics
- 1.3 Semantic Models
- 1.4 The von Neumann Model of Computation
- 1.5 Objects
- 1.6 Sets
- 1.7 Relations
- 1.8 Functions
- 1.9 Problems

Chapter 2 Expressions and Their Notation

- 2.1 Expressions
- 2.2 Function Definitions
- 2.3 Formal Syntax Notation
- 2.4 Problems

Chapter 3 Symbolic Expressions and Abstract Programs

- 3.1 Symbolic Expressions
- 3.2 Abstract Programs
- 3.3 Recursion
- 3.4 Abstract Syntax
- 3.5 Problems

Part II Function-Based Computing

Chapter 4 Lambda Calculus

- 4.1 Syntax of Lambda Calculus
- 4.2 General Model of Computation
- 4.3 Standard Simplification
- 4.4 Identifiers
- 4.5 Substitution Rules

- 4.6 Conversion Rules, Reduction, and Normal Order
- 4.7 The Church-Rosser Theorem
- 4.8 Order of Evaluation
- 4.9 Multiple Arguments, Currying, and Naming Functions
- 4.10 Basic Arithmetic in Lambda Calculus
- 4.11 Boolean Operations in Lambda Calculus
- 4.12 Recursion in Lambda Calculus
- 4.13 Problems

Chapter 5 A Formal Basis for Abstract Programs

- 5.1 A Basic Syntax
- 5.2 Constants
- 5.3 Function Application
- 5.4 Conditional Expressions
- 5.5 Let Expressions-Local Definitions
- 5.6 Recursive Definitions
- 5.7 Global Definitions
- 5.8 Higher-Order Functions
- 5.9 An Example-Symbolic Differentiation
- 5.10 Problems

Chapter 6 Self-Interpretation

- 6.1 Abstract Interpreters
- 6.2 Lambda Expressions as S-Expressions
- 6.3 An Expanded Interpreter
- 6.4 Association Lists
- 6.5 Closures
- 6.6 Recursive Closures
- 6.7 Closing the Loop—Read-Eval-Print
- 6.8 Problems

Chapter 7 The SECD Abstract Machine

- 7.1 List Memory
- 7.2 Basic Data Structures
- 7.3 The SECD Machine Structures
- 7.4 The Basic Instruction Set
- 7.5 Compiling Simple S-expressions
- 7.6 Sample Compilation

7.7 The Funarg Problem

7.8 Optimizations

7.9 Problems

Chapter 8 Memory Management for S-expressions

8.1 Allocation

8.2 Mark-Sweep Collection

8.3 Reference Counts

8.4 Garbage Collection

8.5 Alternative List Representations

8.6 Problems

Chapter 9 Demand-Drive Evaluation

9.1 Explicit Delays and Forces

9.2 Program Examples

9.3 Recipes, Promises, Futures

9.4 Lazy Evaluation

9.5 Streams

9.6 Continuations

9.7 Problems

Chapter 10 LISP: Variations and Implementations

10.1 The Original LISP

10.2 Scheme—A Pure LISP

10.3 Common LISP—A Modern Standard

10.4 MultiLISP—A Parallel LISP

10.5 The CADR Machine

10.6 CADR Derivatives

10.7 Other LISP Machines

10.8 Benchmarking

10.9 Problems

Chapter 11 Combinators and Graph Reduction

11.1 Basic Combinators

11.2 Bracket Abstraction

11.3 Additional Combinators

11.4 Optimizations

11.5 Graph Reduction

11.6 Supercombinators

11.7 Combinator Machines

11.8 The G-Machine

11.9 Problems

Chapter 12 Other Function-Based Computing Systems

12.1 Logo

12.2 FP

12.3 Hope

12.4 Tree Machines

12.5 Alice

12.6 Rediflow

12.7 Function Caching

12.8 Problems

Part III Logic-Based Computing

Chapter 13 A Logic Overview

13.1 Informal Overview

13.2 Formal Logic Systems

13.3 Propositional Logic

13.4 A Simple Inference Engine

13.5 A Simple Problem

13.6 Problems

Chapter 14 Predicate Logic and the First Inference Engine

14.1 Basic Syntax

14.2 Basic Interpretations

14.3 Standard Expression Forms

14.4 Horn Clauses

14.5 Decidability Issues—The Herbrand Result

14.6 The Herbrand Inference Engine

14.7 Problems

Chapter 15 Fundamentals of Practical Inference Engines

15.1 Pattern Matching and Unification

15.2 Common Inference Rules

15.3 Resolution-Based Inference Rules

15.4 Refutation Completeness

15.5 Decision Procedures

15.6	Strategies—An Overview
15.7	Example
15.8	Deduction Trees
15.9	Problems
Chapter 16	The PROLOG Inference Engine
16.1	Structure of a PROLOG Program
16.2	Procedural Interpretation
16.3	A PROLOG-Like Propositional Inference Engine
16.4	The PROLOG Inference Engine
16.5	Special Features of PROLOG
16.6	Some Examples
16.7	Problems
Chapter 17	The Warren Abstract Machine
17.1	Program Structure
17.2	Major Data Structures and State Registers
17.3	Memory Word Format
17.4	Simplified Choice Point
17.5	Simplified WAM Instruction Set
17.6	PROLOG-to-WAM Compiler Overview
17.7	Supporting Built-Ins
17.8	Detailed Example
17.9	Problems
Chapter 18	Optimizations and Extensions
18.1	An Optimized WAM Summary
18.2	Optimizing Single Clauses
18.3	Optimizing Multiple Clauses
18.4	Other Optimizations
18.5	Annotations and Backtrack Control
18.6	Problems
Chapter 19	PROLOG Implementations
19.1	Measuring PROLOG Performance
19.2	Compiling to A Conventional Computer
19.3	The Original Tick and Warren Machine
19.4	Fifth-Generation PROLOG Machines
19.5	The PLM and its Derivatives

- 19.6 The Low RISC Approach
- 19.7 Using Associative Memory
- 19.8 Problems

Chapter 20 All-Solutions Inference Engine

- 20.1 Database Overview
- 20.2 The Relational Model
- 20.3 Computationally Expensive Operators: Joins
- 20.4 Relational Algebra
- 20.5 Implementations
- 20.6 Production-Rule Systems
- 20.7 The Rete Algorithm
- 20.8 An Abstract Machine
- 20.9 Machines
- 20.10 Problems

Chapter 21 Parallel Inference Engines

- 21.1 Taxonomy of Logic Parallelism
- 21.2 Variations of OR Parallelism
- 21.3 Variations of AND Parallelism
- 21.4 AND/OR Process Model
- 21.5 Committed-Choice Languages
- 21.6 Implementations
- 21.7 Problems

References

Index